

# LDA/QDA/KNN

*Peilin Chen*

## R Markdown

This is an R Markdown document. This is an example of using LDA, QDA and KNN for data classification. It is based on one of my homework for statistical learning, written by Robert Tibshirani, springer press. <http://rmarkdown.rstudio.com>.

We use auto as an example dataset for applying LDA

```
require(ISLR)
require(MASS)
require(class)
attach(Auto)
head(Auto)
```

```
##   mpg cylinders displacement horsepower weight acceleration year origin
## 1  18         8          307         130   3504          12.0    70      1
## 2  15         8          350         165   3693          11.5    70      1
## 3  18         8          318         150   3436          11.0    70      1
## 4  16         8          304         150   3433          12.0    70      1
## 5  17         8          302         140   3449          10.5    70      1
## 6  15         8          429         198   4341          10.0    70      1
##                                     name
## 1 chevrolet chevelle malibu
## 2      buick skylark 320
## 3    plymouth satellite
## 4      amc rebel sst
## 5      ford torino
## 6      ford galaxie 500
```

Category the mpg to mpg consuming (1) when the value of mpg is greater than the median mpg, and mpg saving (0)

```
mpg0<-ifelse(Auto$mpg>median(Auto$mpg),1,0)
mydat<-data.frame(Auto,mpg0)
head(mydat)
```

```
##   mpg cylinders displacement horsepower weight acceleration year origin
## 1  18         8          307         130   3504          12.0    70      1
## 2  15         8          350         165   3693          11.5    70      1
## 3  18         8          318         150   3436          11.0    70      1
## 4  16         8          304         150   3433          12.0    70      1
## 5  17         8          302         140   3449          10.5    70      1
## 6  15         8          429         198   4341          10.0    70      1
##                                     name mpg0
## 1 chevrolet chevelle malibu      0
```

```
## 2      buick skylark 320      0
## 3      plymouth satellite    0
## 4      amc rebel sst        0
## 5      ford torino          0
## 6      ford galaxie 500     0
```

separate the dataset to training and testing If the year is even, it is assigned to be training set, if it is odd, it is assigned to be test set.

```
train<-(year%%2==0)
train.auto<-Auto[train,]
test.auto<-Auto[!train,]
mpg0.train<-mpg0[train]
mpg0.test<-mpg0[!train]
```

Applying LDA and get the confusion table for the test dataset

```
lda.auto<-lda(mpg0.train~cylinders+displacement+horsepower+weight,data = train.auto)
lda.pred<-predict(lda.auto,test.auto)
lda.class<-lda.pred$class
Table<-table(lda.class,mpg0.test)
Table
```

```
##      mpg0.test
## lda.class  0  1
##           0 86  9
##           1 14 73
```

The overall test error of the model is 12.64%.

```
type1<-Table[2]/sum(Table[,1])
#type 1 error is false positive, i.e. assume it is mpg saving(0) but turns out to be mpg consuming(1)
type2<-Table[3]/sum(Table[,2])
overall<-mean(lda.class!=mpg0.test)
rbind(type1,type2,overall)
```

```
##      [,1]
## type1 0.1400000
## type2 0.1097561
## overall 0.1263736
```

Applying QDA for classification and get the confusion table in the test dataset. The overall test error is 13.18%

```
qda.fit<-qda(mpg0.train~cylinders+displacement+horsepower+weight, data = train.auto)
qda.pred<-predict(qda.fit,test.auto)
qda.class<-qda.pred$class
Table<-table(qda.class,mpg0.test)
Table
```

```
##          mpg0.test
## qda.class  0  1
##          0 89 13
##          1 11 69
```

Get the type 1 error (False Positive), type 2 error (False Negative rate) and overall error rate on the test dataset

```
type1<-Table[2]/sum(Table[,1])
#type 1 error is false positive, i.e. assume it is mpg saving(0) but turns out to be mpg consuming(1)
type2<-Table[3]/sum(Table[,2])
overall<-mean(qda.class!=mpg0.test)
rbind(type1,type2,overall)
```

```
##          [,1]
## type1    0.1100000
## type2    0.1585366
## overall  0.1318681
```

Using Logistic regression for classification. We calculate the confusion table in the test dataset, and get the Type 1/2 error, overall error rate.

```
logit.fit<-glm(mpg0.train~cylinders+displacement+horsepower+weight, data = train.auto,family =binomial)
summary(logit.fit)
```

```
##
## Call:
## glm(formula = mpg0.train ~ cylinders + displacement + horsepower +
##      weight, family = binomial, data = train.auto)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.48027  -0.03413   0.10583   0.29634   2.57584
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  17.658730   3.409012   5.180 2.22e-07 ***
## cylinders    -1.028032   0.653607  -1.573  0.1158
## displacement  0.002462   0.015030   0.164  0.8699
## horsepower   -0.050611   0.025209  -2.008  0.0447 *
## weight       -0.002922   0.001137  -2.569  0.0102 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 289.58  on 209  degrees of freedom
## Residual deviance:  83.24  on 205  degrees of freedom
## AIC: 93.24
##
## Number of Fisher Scoring iterations: 7
```

```

prob<-predict(logit.fit,test.auto, type = "response")
logit.pred<-rep("0",length(mpg0.test))
logit.pred[prob>0.5]<-"1"
Table<-table(logit.pred, mpg0.test)
Table

```

```

##           mpg0.test
## logit.pred  0  1
##           0 89 11
##           1 11 71

```

```

mean(logit.pred!=mpg0.test)

```

```

## [1] 0.1208791

```

```

type1<-Table[2]/sum(Table[,1])
type2<-Table[3]/sum(Table[,2])
overall<-mean(logit.pred!=mpg0.test)
rbind(type1,type2,overall)

```

```

##           [,1]
## type1    0.1100000
## type2    0.1341463
## overall  0.1208791

```

Knn method, we need to tune the number of neighbors borrowed for calculation. To tune the number of neighbors (n), write a function knnfun() with n as input

```


set.seed(010)
knnfun<-function(n){
  type1<-type2<-r<-NULL
  for (i in 1:n){
    knn.fit<-knn(train.auto[,2:5],test.auto[,2:5], mpg0.train,k=i)
    Table<-table(knn.fit, mpg0.test)
    r[i]<-mean(knn.fit!=mpg0.test)
    type1[i]<-Table[2]/sum(Table[,1])
    type2[i]<-Table[3]/sum(Table[,2])
  }
  plot(1:n,r, xlim = c(1,n),ylim = c(0,0.5),
       xlab = "value of k", ylab="error", type = "l")
  abline(h=min(r), col= "red", lty=2)
  legend("bottom",paste(c("k=", "min overall error"),c(which(r==min(r)),round(min(r),3))),
        bty = "n", cex=0.6)
  par(new=TRUE)
  plot(1:n,type1,type = "l",ylim=c(0,0.5),xaxt="n",yaxt="n",xlab = "",
       ylab="", col="blue",lty=3)
  par(new=TRUE)
  plot(1:n,type2,type="l",ylim = c(0,0.5), xaxt="n",yaxt="n",xlab = "", ylab="",
       col="green", lty=5)
  legend("top",c("type1","type2", "overall"),
        lty = c(3,5,2),col=c("blue","green","black"),

```

```

    bty = "n", cex=0.6)
  return (cbind(type1,type2,r))
}
knnfun(100)

```



	type1	type2	r
## [1,]	0.17	0.13414634	0.1538462
## [2,]	0.13	0.10975610	0.1208791
## [3,]	0.16	0.10975610	0.1373626
## [4,]	0.16	0.10975610	0.1373626
## [5,]	0.18	0.10975610	0.1483516
## [6,]	0.18	0.09756098	0.1428571
## [7,]	0.20	0.08536585	0.1483516
## [8,]	0.20	0.08536585	0.1483516
## [9,]	0.22	0.08536585	0.1593407
## [10,]	0.22	0.08536585	0.1593407
## [11,]	0.21	0.08536585	0.1538462
## [12,]	0.22	0.09756098	0.1648352
## [13,]	0.21	0.08536585	0.1538462
## [14,]	0.18	0.10975610	0.1483516
## [15,]	0.20	0.10975610	0.1593407
## [16,]	0.20	0.09756098	0.1538462
## [17,]	0.21	0.08536585	0.1538462
## [18,]	0.19	0.08536585	0.1428571
## [19,]	0.20	0.08536585	0.1483516
## [20,]	0.19	0.08536585	0.1428571
## [21,]	0.18	0.08536585	0.1373626
## [22,]	0.19	0.09756098	0.1483516
## [23,]	0.18	0.08536585	0.1373626
## [24,]	0.17	0.10975610	0.1428571
## [25,]	0.17	0.10975610	0.1428571

```

## [26,] 0.17 0.10975610 0.1428571
## [27,] 0.17 0.09756098 0.1373626
## [28,] 0.17 0.09756098 0.1373626
## [29,] 0.17 0.09756098 0.1373626
## [30,] 0.17 0.09756098 0.1373626
## [31,] 0.17 0.09756098 0.1373626
## [32,] 0.17 0.08536585 0.1318681
## [33,] 0.18 0.08536585 0.1373626
## [34,] 0.18 0.08536585 0.1373626
## [35,] 0.18 0.09756098 0.1428571
## [36,] 0.18 0.08536585 0.1373626
## [37,] 0.18 0.08536585 0.1373626
## [38,] 0.18 0.08536585 0.1373626
## [39,] 0.18 0.08536585 0.1373626
## [40,] 0.18 0.08536585 0.1373626
## [41,] 0.18 0.08536585 0.1373626
## [42,] 0.18 0.08536585 0.1373626
## [43,] 0.18 0.08536585 0.1373626
## [44,] 0.19 0.09756098 0.1483516
## [45,] 0.18 0.08536585 0.1373626
## [46,] 0.19 0.08536585 0.1428571
## [47,] 0.18 0.08536585 0.1373626
## [48,] 0.19 0.08536585 0.1428571
## [49,] 0.19 0.08536585 0.1428571
## [50,] 0.19 0.08536585 0.1428571
## [51,] 0.19 0.08536585 0.1428571
## [52,] 0.19 0.08536585 0.1428571
## [53,] 0.19 0.08536585 0.1428571
## [54,] 0.20 0.08536585 0.1483516
## [55,] 0.20 0.08536585 0.1483516
## [56,] 0.19 0.08536585 0.1428571
## [57,] 0.19 0.08536585 0.1428571
## [58,] 0.21 0.08536585 0.1538462
## [59,] 0.20 0.08536585 0.1483516
## [60,] 0.20 0.08536585 0.1483516
## [61,] 0.20 0.08536585 0.1483516
## [62,] 0.20 0.08536585 0.1483516
## [63,] 0.20 0.08536585 0.1483516
## [64,] 0.19 0.08536585 0.1428571
## [65,] 0.19 0.08536585 0.1428571
## [66,] 0.19 0.08536585 0.1428571
## [67,] 0.19 0.08536585 0.1428571
## [68,] 0.19 0.08536585 0.1428571
## [69,] 0.19 0.08536585 0.1428571
## [70,] 0.18 0.08536585 0.1373626
## [71,] 0.18 0.08536585 0.1373626
## [72,] 0.19 0.08536585 0.1428571
## [73,] 0.18 0.08536585 0.1373626
## [74,] 0.18 0.09756098 0.1428571
## [75,] 0.18 0.09756098 0.1428571
## [76,] 0.18 0.08536585 0.1373626
## [77,] 0.18 0.08536585 0.1373626
## [78,] 0.19 0.08536585 0.1428571
## [79,] 0.18 0.08536585 0.1373626

```

```
## [80,] 0.17 0.09756098 0.1373626
## [81,] 0.19 0.08536585 0.1428571
## [82,] 0.18 0.09756098 0.1428571
## [83,] 0.19 0.08536585 0.1428571
## [84,] 0.19 0.08536585 0.1428571
## [85,] 0.19 0.08536585 0.1428571
## [86,] 0.19 0.08536585 0.1428571
## [87,] 0.20 0.08536585 0.1483516
## [88,] 0.20 0.08536585 0.1483516
## [89,] 0.20 0.08536585 0.1483516
## [90,] 0.21 0.08536585 0.1538462
## [91,] 0.21 0.08536585 0.1538462
## [92,] 0.19 0.08536585 0.1428571
## [93,] 0.19 0.08536585 0.1428571
## [94,] 0.20 0.08536585 0.1483516
## [95,] 0.19 0.08536585 0.1428571
## [96,] 0.19 0.08536585 0.1428571
## [97,] 0.19 0.08536585 0.1428571
## [98,] 0.19 0.08536585 0.1428571
## [99,] 0.19 0.08536585 0.1428571
## [100,] 0.19 0.08536585 0.1428571
```


If we define the False Positive as  $P(mpg = 1 | mpg = 0)$  and False Negative as  $P(mpg = 0 | mpg = 1)$ . In this case, if we consider False Positive as more severe, and we want to specifically control type 1 error (False Positive rate). Then the ROC curve might help us compare the performance of the classification methods. A good method should have small False Positive rate and high True Positive rate, which leads to the curve approach the left upper corner

```
#Roc curve for lda
ROC_curv2<-function(n){
# changing the cut off of classification
lda.auto <- lda(mpg0.train~cylinders+displacement+horsepower+weight,data = train.auto)
lda.pred<-predict(lda.auto,test.auto)
pr_mpg1<-lda.pred$posterior[,1]
#lda.pred posterior probability [,1] is
#the probability that belongs to mpg=1
truepos <- numeric(n)
falsepos <- numeric(n)
cls<-numeric(length(mpg0.test))
p1 <- (1:n)/(n+1)
for (i in 1:n){
  p <- p1[i]
  cls<-ifelse(pr_mpg1>p,1,0)
  Table<-table(cls,mpg0.test)
  falsepos[i]<-Table[3]/sum(Table[,2])
#type 1 error is false positive, i.e. assume it is mpg saving(0) but turns out to be mpg consuming(1)
  truepos[i]<-1-Table[2]/sum(Table[,1])
}

plot(falsepos ~ truepos, type = "l",
     xlab = "False positive rate",
     ylab = "True positive rate (Sensitivity)",
     main="ROC Curve for testing data based on LDA", cex=0.8)
return(cbind(falsepos,truepos))
```

```
}
result<-ROC_curv2(10000)
```

### ROC Curve for testing data based on LDA



```
#result[,1]
pvalue1<-(which(result[,2]<0.1))/10000
range(pvalue1)
```

```
## [1] 0.0007 0.0530
```

```
pvalue2<-(which(result[,1]>0.8))/10000
range(pvalue2)
```

```
## [1] 0.0306 1.0000
```

similar results can be get for QDA and logistic regression